

Data Analysis: Statistical Modeling and Computation in Applications

Gradient Descent

Outline

- Convex functions
- Gradient descent: main scheme
- Direction
- Step size
- Convergence
- Stochastic Gradient Descent

Recall: Ordinary least Squares Estimator (OLS)

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^N (y_i - \underbrace{\mathbf{x}_i \mathbf{w}}_{\hat{y}_i})^2$$

- Today: general method for finding minima

Notation: to avoid confusion with “typical” notation in optimization vs statistics, for today’s lecture, we replace β by \mathbf{w} : \mathbf{w} are the parameters to optimize.

Recall: Ordinary least Squares Estimator (OLS)

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^N (y_i - \mathbf{x}_i \mathbf{w})^2$$

- Today: general method for finding minima
- Recall:
setting derivative to zero gives *normal equations* and closed form $\hat{\mathbf{w}}$

Notation: to avoid confusion with “typical” notation in optimization vs statistics, for today’s lecture, we replace β by \mathbf{w} : \mathbf{w} are the parameters to optimize.


Recall: Ordinary least Squares Estimator (OLS)

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^N (y_i - \mathbf{x}_i \mathbf{w})^2$$


- Today: general method for finding minima
- Recall:
setting derivative to zero gives *normal equations* and closed form $\hat{\mathbf{w}}$
- *When does setting the derivative to zero give the minimum?*

Notation: to avoid confusion with “typical” notation in optimization vs statistics, for today’s lecture, we replace β by \mathbf{w} : \mathbf{w} are the parameters to optimize.

Intuition




convex



non-convex


Convexity: “bowl-shapedness”



Function f is **convex** if at each point, the gradient gives a linear lower bound, i.e., for all u, w :

$$\underline{f(u)} \geq \underline{f(w) + \langle \nabla f(w), u - w \rangle}. \quad \leftarrow$$

Convexity: “bowl-shapedness”




Function f is **convex** if at each point, the gradient gives a linear lower bound, i.e., for all u, w :

$$f(u) \geq f(w) + \langle \nabla f(w), u - w \rangle.$$

- if $\nabla f(w) = 0$ (local property), then w is a **global minimum** (global):
 $f(u) \geq f(w) + \langle 0, u - w \rangle$ for all u . Not for non-convex functions!

Convexity: “bowl-shapedness”




Function f is **convex** if at each point, the gradient gives a linear lower bound, i.e., for all u, w :

$$f(u) \geq f(w) + \langle \nabla f(w), u - w \rangle.$$

- if $\nabla f(w) = 0$ (local property), then w is a **global minimum** (global):
 $f(u) \geq f(w) + \langle 0, u - w \rangle$ for all u . Not for non-convex functions!
- Hence, for convex f , all we need to find is a point with $\nabla f(w) = 0$.

Convexity: “chord across bowl”



- “chord across bowl”: f is convex if for all u, w and $0 \leq \lambda \leq 1$:

$$f(\underbrace{\lambda w + (1 - \lambda)u}_V) \leq \underbrace{\lambda f(w) + (1 - \lambda)f(u)}_*$$

Convexity: 3 criteria

- ① “linear lower bound”
- ② “chord across bowl”

Convexity: 3 criteria

- ① “linear lower bound”
- ② “chord across bowl”
- ③ if f is twice differentiable: f is convex if for all w
 $\nabla^2 f(w)$ is positive semidefinite (in 1D: $f''(w) \geq 0$).

$$A \text{ is psd if } v^T A v \geq 0 \quad \forall v$$

Examples of convex functions

scalar

$$f(w) = bw$$

vector

$$f(w) = b^\top w$$

for constant b

Examples of convex functions

scalar

$$f(w) = bw$$

$$f(w) = aw^2$$

vector

$$f(w) = b^\top w$$

$$f(w) = w^\top \mathbf{A} w$$

for constant b

for constant $a \geq 0$,

\mathbf{A} positive semidefinite

Examples of convex functions

scalar

$$f(w) = bw$$

$$f(w) = aw^2$$

$$f(w) = \exp(bw)$$

vector

$$f(w) = b^\top w$$

$$f(w) = w^\top \mathbf{A} w$$

$$f(w) = \exp(b^\top w)$$

for constant b

for constant $a \geq 0$,

\mathbf{A} positive semidefinite

for constant b

Examples of convex functions

scalar

$$f(w) = bw$$

$$f(w) = aw^2$$

$$f(w) = \exp(bw)$$

vector

$$f(w) = b^\top w$$

$$f(w) = w^\top \mathbf{A} w$$

$$f(w) = \exp(b^\top w)$$

for constant b

for constant $a \geq 0$,

\mathbf{A} positive semidefinite

for constant b

If $f(w)$ and $g(w)$ are convex functions, then $f(w) + g(w)$ is convex.

\Rightarrow least squares loss is convex!

Outline

- Convex functions
- Gradient descent: main scheme
- Direction
- Step size
- Convergence
- Stochastic Gradient Descent

- sometimes, solving $\nabla f(w) = 0$ is difficult directly:
 - nonlinear equations
 - matrix inversion expensive for large matrices ...

- sometimes, solving $\nabla f(w) = 0$ is difficult directly:
 - nonlinear equations
 - matrix inversion expensive for large matrices ...
- **Optimization:** essential ingredient of modern data science.
find

$$\min_{w \in \mathbb{R}^p} f(w)$$

Examples: minimizing (regularized) loss/error, maximizing likelihood, ...

- sometimes, solving $\nabla f(w) = 0$ is difficult directly:
 - nonlinear equations
 - matrix inversion expensive for large matrices ...
- **Optimization:** essential ingredient of modern data science.
find



$$\min_{w \in \mathbb{R}^p} f(w)$$

Examples: minimizing (regularized) loss/error, maximizing likelihood, ...

- typically iteratively



Optimization: Gradient Descent

$$\min_{w \in \mathbb{R}^p} f(w)$$



Optimization: Gradient Descent

$$\min_{w \in \mathbb{R}^p} f(w)$$





Basic idea: greedy local search

- Start with an arbitrary guess w^0 .
- In each iteration, move in direction of “progress” (determined *locally*)

Optimization: Gradient Descent

$$\min_{w \in \mathbb{R}^p} f(w)$$



Basic idea: greedy local search

- Start with an arbitrary guess w^0 .
- In each iteration, move in direction of “progress” (determined *locally*)
- if done right, finds point w with $\nabla f(w) \approx 0$.
 - convex f : global minimum.
 - non-convex f : local minimum, local maximum or saddle point

General scheme

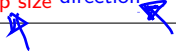
start with some w^0
for $t = 0, 1, 2, \dots$

$$w^{t+1} \leftarrow \underbrace{w^t}_{\text{step size}} + \underbrace{\alpha_t}_{\text{direction}} d^t$$

$$w_i^{t+1} \leftarrow w_i^t + \alpha d_i^t$$

General scheme

start with some w^0
for $t = 0, 1, 2, \dots$


$$w^{t+1} \leftarrow w^t + \underbrace{\alpha_t}_{\text{step size}} \underbrace{d^t}_{\text{direction}}$$


Questions:

- 1 which direction d^t ?
- 2 step size α_t ?
- 3 how many iterations?

Deriving GD: Intuition

Solving $\nabla f(w) = 0$ is difficult in general, but easy for “nice” quadratic functions: Approximate f locally with quadratic function g_t , minimize g_t .



Deriving GD: Quadratic Upper bounds

- **Want:** (1) $g_t(w) \geq f(w) \ \forall w$ and (2) $g(w^t) = f(w^t)$.

Deriving GD: Quadratic Upper bounds

- **Want:** (1) $g_t(w) \geq f(w) \forall w$ and (2) $g(w^t) = f(w^t)$.
- Use *Taylor expansion* to get $g_t(w)$: in 1D

Deriving GD: Quadratic Upper bounds


- **Want:** (1) $g_t(w) \geq f(w) \forall w$ and (2) $g(w^t) = f(w^t)$.
- Use *Taylor expansion* to get $g_t(w)$: in 1D

$$f(u) = \underbrace{f(w^t) + f'(w^t)(u - w^t)} + \underbrace{\frac{1}{2}f''(w^t)(u - w^t)^2} + \dots$$

Deriving GD: Quadratic Upper bounds

- **Want:** (1) $g_t(w) \geq f(w) \forall w$ and (2) $g(w^t) = f(w^t)$.
- Use *Taylor expansion* to get $g_t(w)$: in 1D

$$f(u) = f(w^t) + f'(w^t)(u - w^t) + \frac{1}{2}f''(w^t)(u - w^t)^2 + \dots$$

$$\supseteq g_t(u) = f(w^t) + f'(w^t)(u - w^t) + \frac{L}{2}(u - w^t)^2$$


- Choose L large enough to satisfy (1).

Deriving GD: Quadratic Upper bounds

- **Want:** (1) $g_t(w) \geq f(w) \forall w$ and (2) $g(w^t) = f(w^t)$.
- Use *Taylor expansion* to get $g_t(w)$: in 1D

$$f(u) = f(w^t) + f'(w^t)(u - w^t) + \frac{1}{2}f''(w^t)(u - w^t)^2 + \dots$$

$$g_t(u) = f(w^t) + f'(w^t)(u - w^t) + \frac{L}{2}(u - w^t)^2$$

- Choose L large enough to satisfy (1).
- Setting $g'_t(u) = 0$ gives minimizer of g :

$$u_t = w^t - \frac{1}{L} f'(w^t) = w^{t+1}$$
$$u_t = w^t - \frac{1}{L} \nabla_w f(w^t) = w^{t+1}$$

Deriving GD: Quadratic Upper bounds

- **Want:** (1) $g_t(w) \geq f(w) \forall w$ and (2) $g(w^t) = f(w^t)$.
- Use *Taylor expansion* to get $g_t(w)$: in 1D

$$f(u) = f(w^t) + f'(w^t)(u - w^t) + \frac{1}{2}f''(w^t)(u - w^t)^2 + \dots$$

$$g_t(u) = f(w^t) + f'(w^t)(u - w^t) + \frac{L}{2}(u - w^t)^2$$

- Choose L large enough to satisfy (1).
- Setting $g'_t(u) = 0$ gives minimizer of g :

$$u_t = w^t - \frac{1}{L}f'(w^t)$$

Set $w^{t+1} = u_t = w^t - \frac{1}{L}f'(w^t)$, i.e., $d_t = -f'(w^t)$, $\alpha = 1/L$.

Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.

Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.
- Gradient for squared loss:

$$\nabla_w \left(\sum_{i=1}^N (y_i - x_i \cdot w)^2 \right) = \sum_{i=1}^N \nabla_w (y_i - x_i \cdot w)^2$$

Handwritten blue annotations: Below the first y_i is y_i and below the $-$ is $- \hat{y}_i$. Below the second y_i is \hat{y}_i . Below the ∇_w in the second sum is a double underline.

Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.
- Gradient for squared loss:

$$\begin{aligned}\nabla_w \left(\sum_{i=1}^N (y_i - x_i \cdot w)^2 \right) &= \sum_{i=1}^N \nabla_w (y_i - x_i \cdot w)^2 \\ &= -2 \sum_{i=1}^N \underbrace{(y_i - x_i \cdot w)}_{\gamma_i - \hat{y}_i} x_i^\top\end{aligned}$$

Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.
- Gradient for squared loss:

$$\begin{aligned}\nabla_w \left(\sum_{i=1}^N (y_i - x_i \cdot w)^2 \right) &= \sum_{i=1}^N \nabla_w (y_i - x_i \cdot w)^2 \\ &= -2 \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top\end{aligned}$$

$$\text{so } w^{t+1} = w^t + \underbrace{2\alpha_t \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top}_{\text{blue bracket and arrow pointing to } x_i^\top}$$

Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.
- Gradient for squared loss:

$$\begin{aligned}\nabla_w \left(\sum_{i=1}^N (y_i - x_i \cdot w)^2 \right) &= \sum_{i=1}^N \nabla_w (y_i - x_i \cdot w)^2 \\ &= -2 \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top\end{aligned}$$

so $w^{t+1} = w^t + 2\alpha_t \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top$

- w^{t+1} will be a combination of x_i 's



Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.
- Gradient for squared loss:

$$\begin{aligned}\nabla_w \left(\sum_{i=1}^N (y_i - x_i \cdot w)^2 \right) &= \sum_{i=1}^N \nabla_w (y_i - x_i \cdot w)^2 \\ &= -2 \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top\end{aligned}$$

$$\text{so } w^{t+1} = w^t + 2\alpha_t \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top$$

- w^{t+1} will be a combination of x_i 's
- positive residual ($y_i > \hat{y}_i$):
add fraction of x_i to w^t , increases dot product $x_i \cdot w$

Example: least squares regression

- We obtain iteration: $w^{t+1} \leftarrow w^t - \alpha_t \nabla f(w^t)$.
- Gradient for squared loss:

$$\begin{aligned}\nabla_w \left(\sum_{i=1}^N (y_i - x_i \cdot w)^2 \right) &= \sum_{i=1}^N \nabla_w (y_i - x_i \cdot w)^2 \\ &= -2 \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top\end{aligned}$$

$$\text{so } w^{t+1} = w^t + 2\alpha_t \sum_{i=1}^N (y_i - x_i \cdot w) x_i^\top$$

- w^{t+1} will be a combination of x_i 's
- positive residual ($y_i > \hat{y}_i$):
add fraction of x_i to w^t , increases dot product $x_i \cdot w$
- negative residual: subtract fraction of x_i , decreases dot product

General scheme

start with some w^0
for $t = 0, 1, 2, \dots$


$$w^{t+1} \leftarrow w^t + \underbrace{\alpha_t}_{\text{step size}} \underbrace{d^t}_{\text{direction}}$$

Questions:

- 1 which direction d^t ?
- 2 step size α_t ?
- 3 how many iterations?

More on step sizes

- Step size determined by quadratic upper bound / increase in slope
- If increase is different in different directions: use smallest step size. slower progress overall.







More on step sizes

- Step size determined by quadratic upper bound / increase in slope
- If increase is different in different directions: use smallest step size. slower progress overall.
- If we don't know L : Tuning...

More on step sizes





- Step size determined by quadratic upper bound / increase in slope
- If increase is different in different directions: use smallest step size. slower progress overall.
- If we don't know L : Tuning...
 - too small:** slow progress.
 - too large:** erratic or no convergence.

Step sizes & progress








more demo: http://fa.bianp.net/teaching/2018/eecs227at/gradient_descent.html





Step sizes & progress



Step sizes & progress



Step sizes: using $\alpha = 1/L$



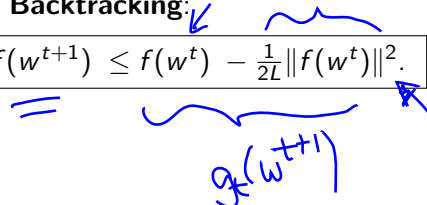
More on step sizes

- Step size determined by quadratic upper bound / increase in slope
- If increase is different in different directions: use smallest step size. slower progress overall.
- If we don't know L : Tuning, or **Backtracking**:

More on step sizes

- Step size determined by quadratic upper bound / increase in slope
- If increase is different in different directions: use smallest step size. slower progress overall.
- If we don't know L : Tuning, or **Backtracking**:

With step size $\alpha_t = \underline{1/L}$: $f(w^{t+1}) \leq f(w^t) - \frac{1}{2L} \|f(w^t)\|^2.$



More on step sizes

- Step size determined by quadratic upper bound / increase in slope
- If increase is different in different directions: use smallest step size. slower progress overall.
- If we don't know L : Tuning, or **Backtracking**:

With step size $\alpha_t = \underline{1/L}$: $f(w^{t+1}) \leq f(w^t) - \frac{1}{2L} \|\nabla f(w^t)\|^2$.

in each step t : select optimistic α_t , check if

$$\|a\|^2 = \sum_{j=1}^n a_j^2$$

$$f(\underline{w_t - \alpha_t \nabla f(w_t)}) \leq f(w_t) - \frac{\alpha_t}{2} \|\underline{\nabla f(w_t)}\|^2$$

If yes: use α_t . If no: $\alpha_t = \underline{\alpha_t/2}$ and check again.

How many iterations?

- **practice:** e.g. until $\|\nabla f(w^t)\|$ is “small enough”, or only small change in loss

$$f(w^t) - f(w^{t+1})$$

¹ means f is also lower bounded by a quadratic, with constant m (instead of L). See e.g. Boyd & Vandenberghe book.


How many iterations?

- **practice:** e.g. until $\|\nabla f(w)\|$ is “small enough”, or only small change in loss
- **theory:** $f(\underline{w^t}) - f(\underline{w^*}) \leq \dots$ (gap to optimum)

¹ means f is also lower bounded by a quadratic, with constant m (instead of L). See e.g. Boyd & Vandenberghe book.

How many iterations?

- **practice:** e.g. until $\|\nabla f(w)\|$ is “small enough”, or only small change in loss
- **theory:** $f(w^t) - f(w^*) \leq \dots$ (gap to optimum)
- *convex functions with bounded L (gradients):*

$$f(w^t) - f(w^*) \leq \frac{L \|w^0 - w^*\|^2}{2t}$$


(need $O(1/\epsilon)$ iterations for $f(w^t) - f(w^*) \leq \underline{\underline{\epsilon}}$)

¹ means f is also lower bounded by a quadratic, with constant m (instead of L). See e.g. Boyd & Vandenberghe book.

How many iterations?

- **practice:** e.g. until $\|\nabla f(w)\|$ is “small enough”, or only small change in loss
- **theory:** $f(w^t) - f(w^*) \leq \dots$ (gap to optimum)
- *convex functions with bounded L (gradients):*

$$f(w^t) - f(w^*) \leq \frac{L\|w^0 - w^*\|^2}{2t}$$

(need $O(1/\epsilon)$ iterations for $f(w^t) - f(w^*) \leq \epsilon$)

- *m-strongly convex functions*¹:

$$f(w^t) - f(w^*) \leq \left(1 - \frac{m}{L}\right)^t (f(w^0) - f(w^*))$$

(need $O(\log(1/\epsilon))$ iterations for $f(w^t) - f(w^*) \leq \epsilon$)



¹ means f is also lower bounded by a quadratic, with constant m (instead of L). See e.g. Boyd & Vandenberghe book.

Outline

- Convex functions
- Gradient descent: main scheme
- Direction
- Step size
- Convergence
- Stochastic Gradient Descent

What if N is large?

Stochastic Gradient descent

- Setup: f is a sum: $f(w) = \sum_{i=1}^N f_i(w)$
- Gradient: $\nabla_w f(w) = \sum_{i=1}^N \nabla_w f_i(w)$

What if N is large?

Stochastic Gradient descent

- Setup: f is a sum: $f(w) = \sum_{i=1}^N f_i(w)$
- Gradient: $\nabla_w f(w) = \sum_{i=1}^N \nabla_w f_i(w)$
- Idea: estimate sum via few (one) term(s)!

What if N is large?

Stochastic Gradient descent

- Setup: f is a sum: $f(w) = \sum_{i=1}^N f_i(w)$
- Gradient: $\nabla_w f(w) = \sum_{i=1}^N \nabla_w f_i(w)$
- Idea: estimate sum via few (one) term(s)!

start with some w^0

for $t = 0, 1, 2, \dots$

draw (data point) i uniformly at random, $1 \leq i \leq N$

$$w^{t+1} \leftarrow w^t - \alpha_t \nabla f_i(w^t)$$

What if N is large?

Stochastic Gradient descent

- Setup: f is a sum: $f(w) = \sum_{i=1}^N f_i(w)$
- Gradient: $\nabla_w f(w) = \sum_{i=1}^N \nabla_w f_i(w)$
- Idea: estimate sum via few (one) term(s)!

start with some w^0

for $t = 0, 1, 2, \dots$

draw (data point) i uniformly at random, $1 \leq i \leq N$

$$w^{t+1} \leftarrow w^t - \alpha_t \nabla f_i(w^t)$$

- with “right” step size, also converges to minimum. Step size must shrink ($\approx \frac{1}{t+1}$)

What if N is large?

Stochastic Gradient descent

- Setup: f is a sum: $f(w) = \sum_{i=1}^N f_i(w)$
- Gradient: $\nabla_w f(w) = \sum_{i=1}^N \nabla_w f_i(w)$
- Idea: estimate sum via few (one) term(s)!

start with some w^0

for $t = 0, 1, 2, \dots$

draw (data point) i uniformly at random, $1 \leq i \leq N$





$$w^{t+1} \leftarrow w^t - \alpha_t \nabla f_i(w^t)$$

- with “right” step size, also converges to minimum. Step size must shrink ($\approx \frac{1}{t+1}$)
- more erratic, but standard for large data

- Fit a function to observations from $y = \beta_1 x + \beta_2 x^2$ (upper left: observations) via least squares regression with gradient descent and stochastic gradient descent.

- Fit a function to observations from $y = \beta_1 x + \beta_2 x^2$ (upper left: observations) via least squares regression with gradient descent and stochastic gradient descent.
- Bottom row: Loss function (sum of squared residuals) values for gradient descent (GD, left) and stochastic gradient descent (SGD, right). Note that SGD only uses *one* data point per iteration, whereas GD uses all data points in each iteration.

Demo: plots



Summary: Gradient descent

- (Stochastic) Gradient Descent is a pillar of modern machine learning
- iterative method to find a point with zero gradient
- small steps in direction of negative gradient
- for convex functions, finds the global minimum

Convex Optimization

- S. Boyd & L. Vandenberghe. *Convex Optimization*. Available online: <http://stanford.edu/~boyd/cvxbook/>
Parts of Chapters 3.1, (3.2 for additional optional reading); parts of Chapter 9.1 and 9.3

For some background in linear algebra

- very short: Appendix in Boyd & Vandenberghe. Or: Many statistics books have a chapter on it, e.g., D. Freedman. *Statistical Models – Theory and Practice*.
- a bit longer:
T.A. Garrity. *All the Mathematics you missed: But need to know for Graduate School*. Cambridge University Press.

Appendix: some linear algebra concepts

- **rank** of matrix: max. number of linearly independent columns (rows)
- $\mathbf{X}^\top \mathbf{X}$ is invertible if it has full rank (no zero eigenvalues), i.e., if $N \geq p$ and all columns of \mathbf{X} are linearly independent.

Appendix: some linear algebra concepts

- **rank** of matrix: max. number of linearly independent columns (rows)
- **eigenvalue**: an eigenvector \mathbf{u} of a matrix \mathbf{A} and the corresponding eigenvalue λ satisfy $\mathbf{A}\mathbf{u} = \lambda\mathbf{u}$.

$$\lambda_{\max} = \max_{\mathbf{u} \in \mathbb{R}^p, \|\mathbf{u}\|=1} \mathbf{u}^T \mathbf{A} \mathbf{u} \quad \lambda_{\min} = \min_{\mathbf{u} \in \mathbb{R}^p, \|\mathbf{u}\|=1} \mathbf{u}^T \mathbf{A} \mathbf{u}$$

- $\mathbf{X}^T \mathbf{X}$ is invertible if it has full rank (no zero eigenvalues), i.e., if $N \geq p$ and all columns of \mathbf{X} are linearly independent.

Appendix: some linear algebra concepts

- **rank** of matrix: max. number of linearly independent columns (rows)
- **eigenvalue**: an eigenvector \mathbf{u} of a matrix \mathbf{A} and the corresponding eigenvalue λ satisfy $\mathbf{A}\mathbf{u} = \lambda\mathbf{u}$.

$$\lambda_{\max} = \max_{\mathbf{u} \in \mathbb{R}^p, \|\mathbf{u}\|=1} \mathbf{u}^T \mathbf{A} \mathbf{u} \quad \lambda_{\min} = \min_{\mathbf{u} \in \mathbb{R}^p, \|\mathbf{u}\|=1} \mathbf{u}^T \mathbf{A} \mathbf{u}$$

- symmetric \mathbf{A} is *positive semidefinite* (psd) if

$$\mathbf{v}^T \mathbf{A} \mathbf{v} \geq 0 \text{ for all } \mathbf{v} \in \mathbb{R}^p$$

Equivalent: \mathbf{A} is symmetric and all eigenvalues are nonnegative.

- $\mathbf{X}^T \mathbf{X}$ is invertible if it has full rank (no zero eigenvalues), i.e., if $N \geq p$ and all columns of \mathbf{X} are linearly independent.

Appendix: Additional note: What is L ?

- We want (scalar version):

$$f(u) \leq f(w) + f'(w)(u - w) + \frac{L}{2}(u - w)^2$$

Taylor expansion:

$$f(u) = f(w) + f'(w)(u - w) + \frac{1}{2}f''(z)(u - w)^2$$

so $L = \max_z f''(z)$.

Appendix: Additional note: What is L ?

- We want (scalar version):

$$f(u) \leq f(w) + f'(w)(u - w) + \frac{L}{2}(u - w)^2$$

Taylor expansion:

$$f(u) = f(w) + f'(w)(u - w) + \frac{1}{2}f''(z)(u - w)^2$$

so $L = \max_z f''(z)$.

- for vectors, this condition becomes:

$$v^\top [\nabla^2 f(z)] v \leq L v^\top v \quad \text{for } v = (u - w)$$

i.e., L must be larger than largest eigenvalue of $\nabla^2 f(z)$ for all z .

Appendix: Additional note: What is L ?

- We want (scalar version):

$$f(u) \leq f(w) + f'(w)(u - w) + \frac{L}{2}(u - w)^2$$

Taylor expansion:

$$f(u) = f(w) + f'(w)(u - w) + \frac{1}{2}f''(z)(u - w)^2$$

so $L = \max_z f''(z)$.

- for vectors, this condition becomes:

$$v^\top [\nabla^2 f(z)] v \leq L v^\top v \quad \text{for } v = (u - w)$$

i.e., L must be larger than largest eigenvalue of $\nabla^2 f(z)$ for all z .

- if $f(w) = w^\top \mathbf{A}w + bw$, then $\nabla^2 f(z) = \mathbf{A}$ and L is maximum eigenvalue of \mathbf{A} : $L = \lambda_{\max}$